

Clustering Knowledge Articles for AI-Readiness Assessment

An AI-Readiness Evaluation Framework for Knowledge Articles Using Unsupervised Learning
Aligned with Retrieval-Augmented Generation Requirements

Aicel Mae Alvarez
Author

Postgraduate Diploma in Artificial Intelligence and Machine Learning
Asian Institute of Management

Executive Summary

This capstone project delivers an end-to-end, governance-aligned framework to evaluate and cluster enterprise knowledge articles based on their readiness for ingestion by enterprise AI assistants. Using unsupervised learning, natural language processing, and explainability techniques, the project demonstrates how organizations can operationalize AI governance principles through scalable, data-driven content assessment.

The solution translates abstract AI governance requirements such as transparency, quality assurance, risk segmentation, and human oversight into measurable indicators applied at scale. The outcome is a practical decision support capability that classifies knowledge articles into actionable AI-readiness clusters, enabling controlled onboarding of content into AI systems while reducing operational, reputational, and compliance risk.

Designed with enterprise deployment and MLOps considerations in mind, the project aligns academic rigor with real-world AI governance, lifecycle management, and continuous improvement expectations.

Problem Context and Business Relevance

As organizations adopt enterprise AI assistants, knowledge content becomes a governed data asset rather than a static repository. Poorly structured, outdated, or ambiguous articles introduce operational risk, degrade AI output quality, and undermine user trust. Manual content reviews do not scale, lack consistency, and are misaligned with enterprise AI governance models.

This project addresses the following governance-driven business question: **How can large volumes of knowledge articles be systematically assessed, segmented, and governed according to their suitability for enterprise AI consumption?**

From an enterprise AI and risk-management perspective, the solution enables:

- Objective, repeatable assessment of AI ingestion readiness
- Risk-based segmentation of content prior to AI exposure
- Prioritized remediation aligned to governance thresholds
- Auditability and traceability of content quality decisions
- Reduced dependency on subjective or manual audits

Data and Feature Strategy

The dataset consists of publicly available textual knowledge-style content, used as a proxy for enterprise knowledge articles. Each document is transformed into a structured representation using features explicitly mapped to enterprise AI-readiness and governance dimensions.

Feature engineering focuses on signals that affect AI system behavior, including:

- Structural integrity and formatting consistency
- Linguistic clarity, ambiguity, and cognitive load
- Metadata completeness and lifecycle indicators
- Content density, redundancy, and retrieval suitability

By combining semantic embeddings with governance-oriented numerical features, the representation ensures that clustering reflects **AI consumption risk and readiness**, rather than superficial topic similarity. This mirrors enterprise practices where content quality, not subject matter, determines AI ingestion eligibility..

Methodology Overview

The project follows a disciplined, governance-aware data science lifecycle aligned with enterprise AI standards:

1. **Governance-aligned problem framing** with explicit success and risk criteria
2. **Data preprocessing and normalization** to support fairness and comparability
3. **Dimensionality reduction** to improve explainability and operational usability
4. **Multi-model clustering evaluation** to avoid algorithmic lock-in

5. **Model selection** based on quantitative metrics and governance interpretability
6. **Explainability analysis** to support transparency and auditability
7. **Ethical and bias review** to identify unintended exclusion or over-penalization

This methodology reflects how AI systems are increasingly evaluated in regulated and enterprise environments, where technical performance must align with governance, accountability, and trust.

Model Evaluation and Selection

Multiple clustering algorithms were evaluated to balance performance, stability, and interpretability. Model selection was guided by:

- Internal validation metrics
- Cluster cohesion and separation
- Business interpretability of resulting segments

The selected model produced well-defined clusters that could be meaningfully labeled along an AI-readiness spectrum, rather than abstract mathematical groupings.

Key Results and Cluster Interpretation

The final clustering resulted in clearly differentiated article segments that align with enterprise AI governance use cases:

1. **High AI-Readiness:** Articles meeting structural, linguistic, and consistency thresholds suitable for controlled AI ingestion
2. **Moderate AI-Readiness:** Articles usable with defined remediation actions and governance review
3. **Low AI-Readiness:** High-risk content requiring significant restructuring, validation, or deprecation

Explainability techniques were applied to identify the dominant drivers of cluster assignment, enabling transparent justification of decisions and supporting governance workflows such as remediation prioritization, exception handling, and audit review.

Ethical AI and Governance Considerations

An explicit ethical and governance review was conducted to evaluate alignment with responsible AI principles, including fairness, transparency, and human oversight.

Key considerations included:

- Potential bias introduced by language complexity or writing style
- Risk of disadvantaging non-native language authors
- Over-reliance on automated signals for governance decisions

The findings reinforce the principle that automated clustering should function as a **decision support and risk triage mechanism**, not an autonomous gatekeeper. Human review, policy oversight, and documented exception handling remain essential components of responsible enterprise AI deployment.

Deployment and MLOps Readiness

While this project focuses on analytical development through exploratory analysis, feature engineering, model training, and evaluation (Steps 1–5), deployment was intentionally scoped out of the implementation phase. Instead, the notebook establishes deployment and MLOps readiness by design, ensuring that the analytical outputs can be operationalized without rework.

Key design considerations demonstrated in the notebook include:

- Deterministic and reproducible preprocessing steps
- Clear separation of feature engineering, modeling, and evaluation logic
- Stable feature definitions aligned to governance objectives
- Interpretable outputs suitable for downstream workflow integration

These elements align with enterprise AI lifecycle expectations, where proof-of-value models are first validated analytically before progressing to production deployment, automation, and monitoring stages.

Limitations

Key limitations include:

- Use of public data as a proxy for proprietary enterprise content
- Lack of human-validated labels for cluster ground truth
- Dependence on engineered proxies for some readiness criteria

These limitations are acknowledged and addressed through recommendations for future enhancements.

Future Enhancements

Recommended next steps include:

- Incorporating human review scores to semi-supervise clustering
- Integrating generative AI to auto-rewrite low-readiness articles
- Building dashboards for operational tracking
- Expanding evaluation across multilingual content

Conclusion

This capstone demonstrates how unsupervised learning can be operationalized within an enterprise AI governance context to address a critical readiness challenge: ensuring that knowledge assets are suitable, safe, and effective for AI consumption.

By translating governance principles into measurable signals and explainable clusters, the project bridges academic AI methods with real-world enterprise requirements. The resulting framework supports responsible AI adoption by enabling risk-based content onboarding, transparent decision-making, and continuous quality improvement which are key foundations for scalable and trustworthy enterprise AI systems.